

Efficient Computer Vision Models for Silkworm Feeding Prediction and Habitat Analysis

Master's Degree in Artificial Intelligence and Robotics

Venneri Roberta (2019902)
Ruggeri Riccardo (1613249)



SAPIENZA
UNIVERSITÀ DI ROMA



Outline

- ▶ Problem Statement
- ▶ State Of The Art
- ▶ Proposed Method
- ▶ Dataset
- ▶ Experimental Setup
- ▶ Model Evaluation
- ▶ Conclusions
- ▶ References




PROBLEM STATEMENT

- Ensuring efficiency in silkworm rearing is a key factor for sustainable silk production, and it critically depends on accurately tracking feeding conditions.
- There are 2 current issues:
 - **manual monitoring**, which is slow, expensive and subject to human errors
 - **hard to distinguish** worms, leaves, and background

TWO CHALLENGES



Food status **Classification**



Automatic and
Unsupervised
Segmentation of
elements in the image



STATE OF THE ART

- Current problems with the manual monitoring of silkworms has led to the need for automated visual analysis.
- However, traditional CNNs are effective, but capture only local spatial features.
- Transformers are effective but also heavy and difficult to implement.
- Hybrid Models like CNN + transformers improve accuracy but need large datasets and a complex training.



PROPOSED METHOD

1. BINARY CLASSIFICATION

GOAL: determine whether the silkworms need feeding.

- Application of a binary classification on images of silkworm farms, where each image is labelled as “feeding” or “no feeding”.
- The model chosen is a convolutional network known as **MobileNet_v2**, which due to its efficiency in embedded environments is particularly suitable for this task.
- It was decided to use a version of the model pre-trained on ImageNet using the torchvision library, so as to take advantage of general features already learned.
- The last layer was then replaced by a fully connected layer so as to adapt the network to our case, i.e. “feeding” or “no feeding”.



PROPOSED METHOD

2. UNSUPERVISED SEGMENTATION

GOAL: Implement unsupervised methods to segment silkworms, mulberry leaves, and backgrounds without ground truth.

It is possible to divide the methodology followed into 3 steps:

1. Feature Extraction with MobileViT

- Implementation of the **MobileViT** model, which exploits the combination of a CNN and Transformer network applied to feature map patches.
- From the model, features are extracted at an intermediate stage (stage2) and made ready for step 2.



PROPOSED METHOD

2. UNSUPERVISED SEGMENTATION

GOAL: Implement unsupervised methods to segment silkworms, mulberry leaves, and backgrounds without ground truth.

2. Feature clustering

- Features extracted from all images are then randomly sampled and dimensionally reduced with PCA.
- This is where segmentation takes place. Two methods are proposed: **K-means** and Gaussian Mixture Model (**GMM**).



PROPOSED METHOD

2. UNSUPERVISED SEGMENTATION

GOAL: Implement unsupervised methods to segment silkworms, mulberry leaves, and backgrounds without ground truth.

3. Local Segmentation + Augment with Sobel filter

- Finally, the proposed method is used on some images to test its effectiveness. A Sobel filter is also applied to highlight the edges more clearly.
- Features and edges are then concatenated to obtain richer features.
- After repeating PCA and clustering, the results obtained are analysed.



DATASET

- The Dataset used is the **Silkwork rearing data**, provided by the professor in .zip format.
- It contains 1352 images in .jpg format of silkworms, mulberry leaves and background.
- The dataset is loaded via a CSV with the names of the images and their labels, applying resize transformations and conversion to PyTorch tensors.
- A train/validation split (80%-20%) is performed for supervised training, while for unsupervised clustering the entire dataset is used without shuffling.



EXPERIMENTAL SETUP

Dataset

- Silkworm Feeding Dataset (~1300 images)
- Train/Validation split: 80 % / 20 % ($\approx 2,161$ / 541 samples)
- Resized to 224×224 pixels

Preprocessing

- `transforms.Resize((224,224))`
- `transforms.ToTensor()`
- ImageNet normalization (mean & std)



EXPERIMENTAL SETUP

Model Architecture

- MobileNet_v2 pre-trained on ImageNet for binary classification
- Final classification layer replaced with Linear(1280 \rightarrow 2)
- MobileViT for unsupervised clustering

Training

- Optimizer: Adam (lr = $1e-4$, weight_decay = $1e-4$)
- Loss: Cross-Entropy
- Batch size: 16
- Max epochs: 10 with **early stopping** (patience = 2 on val loss)



EXPERIMENTAL SETUP

Pruning & Fine-Tuning

- Global unstructured L_1 pruning on all Conv2d and Linear weights
- Pruning levels tested: 0 %, 10 %, 20 %
- Post-pruning fine-tuning: 3 epochs at reduced lr ($1e-5$)

Metrics & Environment

- Metrics: validation loss, accuracy, “Prune Amount vs. Accuracy” curve
- Hardware: Google Colab GPU (Tesla T4), inference time ≈ 15 s per epoch on CPU



MODEL EVALUATION

Accuracy & Loss

- 97.2 % peak val accuracy, ~ 0.10 val loss
- Early stopping prevented overfitting

Pruning Analysis

- 10 % sparsity \rightarrow accuracy \uparrow to 95.9 %, inference time \downarrow
- 20 % sparsity \rightarrow accuracy drop

Unsupervised Segmentation



- **GMM vs. K-Means:** cleaner class masks with GMM
- **Clustering metrics:** Silhouette score & Davies–Bouldin index used to quantify segmentation quality

Visualization

- Training/validation curves
- Accuracy vs. pruning level
- Silhouette vs. number of clusters



MODEL EVALUATION






MODEL EVALUATION



Prune 0% → Val Acc: 94.46%
Prune 10% → Val Acc: 95.94%
Prune 20% → Val Acc: 95.20%






MODEL EVALUATION


Original



KMeans Segmentation



GMM Segmentation





CONCLUSION

Strong Baseline

- MobileNet-V2 → 97.2 % val accuracy, smooth train/val curves

Effective Pruning

- 10 % sparsity → accuracy ↑ to 95.9 %, inference time ↓
- 20 % sparsity → clear performance drop

Segmentation Findings

- GMM delivers higher Silhouette and lower Davies–Bouldin scores vs K-Means
- Cleaner feeding masks, with some over-segmentation remaining

Next Steps

- Integrate segmentation maps into the classifier
- Apply 8-bit quantization & distillation
- Deploy and benchmark on edge hardware



REFERENCES

- Sandler, M., Howard, A., Zhu, M., Zhmoginov, A., & Chen, L.-C. (2018). *MobileNetV2: Inverted Residuals and Linear Bottlenecks*. CVPR.
- Kingma, D. P., & Ba, J. (2015). *Adam: A Method for Stochastic Optimization*. ICLR.
- Frankle, J., & Carbin, M. (2019). *The Lottery Ticket Hypothesis: Finding Sparse, Trainable Neural Networks*. ICLR.
- McLachlan, G., & Peel, D. (2000). *Finite Mixture Models*. Wiley.
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., ... & Duchesnay, É. (2011). *Scikit-learn: Machine Learning in Python*. JMLR, 12, 2825–2830.
- Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., ... & Chintala, S. (2019). *PyTorch: An Imperative Style, High-Performance Deep Learning Library*. NeurIPS.
- Bishop, C. M. (2006). *Pattern Recognition and Machine Learning*. Springer. (for GMM foundations)
- Rehurek, R., & Sojka, P. (2010). *Software Framework for Topic Modelling with Large Corpora*. LREC. (for K-Means)



**THANK YOU FOR YOUR
ATTENTION!**